

Logistic Regression

전 재 욱

**Embedded System 연구실
성균관대학교**

Outline

- Classification
- Hypothesis representation
- Decision boundary
- Cost function
- Derivation of cost function via MLE
- Simplified cost function and gradient descent
- Advanced optimization
- Multi-class classification

Outline

- Classification
- Hypothesis representation
- Decision boundary
- Cost function
- Derivation of cost function via MLE
- Simplified cost function and gradient descent
- Advanced optimization
- Multi-class classification

Classification

■ Email

- Spam / Not Spam?

■ Online Transactions

- Fraudulent (Yes / No)?

■ Tumor


- Malignant / Benign ?

■ $y \in \{0, 1\}$

- 0: "Negative Class" (e.g., benign tumor)
- 1: "Positive Class" (e.g., malignant tumor)

■ $y \in \{0, 1, 2, 3\}$

Classification



■ Threshold classifier output $h_{\theta}(x)$ at 0.5:

- If $h_{\theta}(x) \geq 0.5$, predict “ $y = 1$ ”
- If $h_{\theta}(x) < 0.5$, predict “ $y = 0$ ”

Classification vs Logistic Regression

- Classification: $y = 0$ or 1

- If we were to use a linear regression for classification,
 $h_{\theta}(x)$ can be > 1 or < 0

- Logistic Regression

- Always $0 \leq h_{\theta}(x) \leq 1$

- One classification algorithm

Outline

- Classification
- Hypothesis representation
- Decision boundary
- Cost function
- Derivation of cost function via MLE
- Simplified cost function and gradient descent
- Advanced optimization
- Multi-class classification

Logistic Regression Model

- Takes a probabilistic approach
to learning discriminative functions (i.e., a classifier)
- $h_{\theta}(x)$ should give $p(y = 1|x; \theta)$
 - Want $0 \leq h_{\theta}(x) \leq 1$

Logistic Regression Model


■ Want $0 \leq h_{\theta}(x) \leq 1$

■
$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

■
$$g(z) = \frac{1}{1 + \exp(-z)}$$

■ Sigmoid function

■ Logistic function



Interpretation of Hypothesis Output

■ $h_{\theta}(x)$

■ Estimated probability that $y = 1$ on input x

■ Example

➤ If $x = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 1 \\ tumor_size \end{bmatrix}$ and $h_{\theta}(x) = 0.7$

tell patient that 70% chance of tumor being malignant

■ $h_{\theta}(x) = P(y = 1|x; \theta)$

■ “probability that $y = 1$, given x , parameterized by θ ”

■ $y = 0$ or 1

■ $P(y = 0|x; \theta) + P(y = 1|x; \theta) = 1$

■ $P(y = 0|x; \theta) = 1 - P(y = 1|x; \theta)$

Another Interpretation

■ $h_{\theta}(x)$

■ Estimated probability that $y = 1$ on input x

Odds of $y = 1$

$$\log \frac{p(y = 1|x; \theta)}{p(y = 0|x; \theta)} = \log \frac{h}{1 - h} = \log \frac{1}{\frac{\exp(-\theta^T x)}{1 + \exp(-\theta^T x)}} \\ = \log \frac{1}{\exp(-\theta^T x)} = \theta^T x = \theta_0 + \theta_1 x_1 + \dots + \theta_n x_n$$

■ Logistic regression assumes that

■ the {log odds} is a linear function of x

Outline

- Classification
- Hypothesis representation
- Decision boundary
- Cost function
- Derivation of cost function via MLE
- Simplified cost function and gradient descent
- Advanced optimization
- Multi-class classification

Logistic Regression

■ $h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + \exp(-\theta^T x)}$


■ Suppose

■ predict “ $y = 1$ ” if $h_{\theta}(x) \geq 0.5$


■ $\theta^T x \geq 0$

■ predict “ $y = 0$ ” if $h_{\theta}(x) < 0.5$


■ $\theta^T x < 0$



Decision Boundary



Decision Boundary



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

Decision boundary

$$\theta_0 = -3, \theta_1 = 1, \theta_2 = 1$$

■ Predict “ $y = 1$ ” if $-3 + x_1 + x_2 \geq 0$

■ i.e. if $x_1 + x_2 \geq 3$

■ $h_{\theta}(x) = 0.5$ when $x_1 + x_2 = 3$



■ Predict “ $y = 0$ ” if $x_1 + x_2 < 3$

Nonlinear Decision Boundary


- Can apply basis function expansion to features, same as with linear regression

$$x = \begin{bmatrix} 1 \\ x_1 \\ \dots \\ x_n \end{bmatrix} \rightarrow \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ x_1 x_2 \\ x_1^2 \\ x_2^2 \\ x_1^2 x_2 \\ x_1 x_2^2 \\ \dots \\ x_n \end{bmatrix}$$

Nonlinear Decision Boundary



Nonlinear Decision Boundary




■ $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$

■ Predict “ $y = 1$ ” if $-1 + x_1^2 + x_2^2 \geq 0$

■ $\theta_0 = -1, \theta_1 = \theta_2 = 0, \theta_3 = \theta_4 = 1$

■ $\theta = \begin{bmatrix} -1 \\ 0 \\ 0 \\ 1 \\ 1 \end{bmatrix}$



■ $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_1^2 x_2 + \theta_5 x_1^2 x_2^2 + \theta_6 x_1^3 x_2 + \dots)$

Outline

- Classification
- Hypothesis representation
- Decision boundary
- **Cost function**
- Derivation of cost function via MLE
- Simplified cost function and gradient descent
- Advanced optimization
- Multi-class classification

Parameter θ

■ Training set

■ m examples: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

■ For each example, $x^{(i)} = \begin{bmatrix} x_0^{(i)} \\ x_1^{(i)} \\ \vdots \\ x_n^{(i)} \end{bmatrix} \in R^{n+1}, \quad x_0^{(i)} = 1,$
 $y \in \{0, 1\}$

■ $h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + \exp(-\theta^T x)}$

■ How to choose parameters θ ?

Cost Function

■ Linear regression

- $J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$

■ Define

- $Cost(h_{\theta}(x^{(i)}), y^{(i)}) = \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$

- $J(\theta) = \frac{1}{m} \sum_{i=1}^m Cost(h_{\theta}(x^{(i)}), y^{(i)})$

■ Then

- $J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2 = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} \left(\frac{1}{1 + \exp(-\theta^T x^{(i)})} - y^{(i)} \right)^2$

Cost Function


■ Then

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2 = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} \left(\frac{1}{1 + \exp(-\theta^T x^{(i)})} - y^{(i)} \right)^2$$

■ Non-convex function


➤ Many local min

➔ Gradient descent may not find the global min



Cost Function


- If $J(\theta)$ is convex
 - Gradient descent can converge the global minimum



Convex Logistic Regression Cost Function

■ Convex logistic regression

$$\blacksquare \text{ Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$



■ $\text{Cost} = 0$ if $y = 1$ and $h_{\theta}(x) = P(y = 1|x; \theta) = 1$


■ $\text{Cost} \rightarrow \infty$ as $h_{\theta}(x) \rightarrow 0$

■ Captures intuition that

➤ if $h_{\theta}(x) = 0$ (predict $P(y = 1|x; \theta) = 0$) but $y = 1$
this learning algorithm will be penalized
by a very large cost

Convex Logistic Regression Cost Function

$$\blacksquare \text{ Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$



■ $\text{Cost} = 0$ if $y = 0$ and $h_{\theta}(x) = P(y = 1|x; \theta) = 0$

■ $\text{Cost} \rightarrow \infty$ as $h_{\theta}(x) \rightarrow 1$

■ Captures intuition that

➤ if $h_{\theta}(x) = 1$ (predict $P(y = 1|x; \theta) = 1$) but $y = 0$
this learning algorithm will be penalized
by a very large cost

Convex Logistic Regression Cost Function

$$Cost(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

■ Always $y = 0$ or 1

■ So, for each training example $(x^{(i)}, y^{(i)})$,

$$■ Cost(h_{\theta}(x^{(i)}), y^{(i)}) = -y^{(i)} \log(h_{\theta}(x)) - (1 - y^{(i)}) \log(1 - h_{\theta}(x))$$

■ For all training examples $(x^{(i)}, y^{(i)})$ for $i = 1, 2, \dots, m$

$$\begin{aligned} ■ J(\theta) &= \frac{1}{m} \sum_{i=1}^m Cost(h_{\theta}(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} [\sum_{i=1}^m \{ y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \}] \end{aligned}$$

Convex Logistic Regression Cost Function

- Logistic regression cost function

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \left[\sum_{i=1}^m \{ y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \} \right] \end{aligned}$$

- To fit parameters θ

- $\min_{\theta} J(\theta)$

- To make a prediction, given new x_{new} :

- $h_{\theta}(x_{new}) = \frac{1}{1 + \exp(-\theta^T x_{new})} = P(y = 1 | x_{new}; \theta)$

Outline

- Classification
- Hypothesis representation
- Decision boundary
- Cost function
- Derivation of cost function via MLE
- Simplified cost function and gradient descent
- Advanced optimization
- Multi-class classification

Derivation of Cost Function via MLE

■ $h_{\theta}(x)$

■ Estimated probability that $y = 1$ on input x

■ $p(y = 1|x; \theta) = h_{\theta}(x), p(y = 0|x; \theta) = 1 - h_{\theta}(x)$

➤ $p(y|x; \theta) = (h_{\theta}(x))^y (1 - h_{\theta}(x))^{1-y}$

■ Given m (independent) training examples,
the likelihood of the parameters, $L(\theta)$,

$$L(\theta) = \prod_{i=1}^m p(y^{(i)}|x^{(i)}; \theta) = \prod_{i=1}^m \left(h_{\theta}(x^{(i)}) \right)^{y^{(i)}} \left(1 - h_{\theta}(x^{(i)}) \right)^{1-y^{(i)}}$$

Derivation of Cost Function via MLE

■ θ_{MLE}

■ which maximizes the likelihood (Maximum likelihood estimation)

$$\begin{aligned}\theta_{MLE} &= \arg \max_{\theta} L(\theta) = \arg \max_{\theta} \prod_{i=1}^m p(y^{(i)} | x^{(i)}; \theta) \\ &= \arg \max_{\theta} \prod_{i=1}^m \left(h_{\theta}(x^{(i)}) \right)^{y^{(i)}} \left(1 - h_{\theta}(x^{(i)}) \right)^{1-y^{(i)}}\end{aligned}$$

Derivation of Cost Function via MLE

■ θ_{MLE}

■ which maximizes the likelihood (Maximum likelihood estimation)

Or

$$\begin{aligned}\theta_{MLE} &= \arg \max_{\theta} \log L(\theta) = \arg \max_{\theta} \log \prod_{i=1}^m p(y^{(i)} | x^{(i)}; \theta) \\ &= \arg \max_{\theta} \sum_{i=1}^m \log p(y^{(i)} | x^{(i)}; \theta) \\ &= \arg \max_{\theta} \sum_{i=1}^m \log \left(h_{\theta}(x^{(i)}) \right)^{y^{(i)}} \left(1 - h_{\theta}(x^{(i)}) \right)^{1-y^{(i)}} \\ &= \arg \max_{\theta} \sum_{i=1}^m \{ y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \}\end{aligned}$$

Derivation of Cost Function via MLE

■ θ_{MLE}

■ which maximizes the likelihood (Maximum likelihood estimation)

$$\theta_{MLE} = \underset{\theta}{\operatorname{argmax}} \sum_{i=1}^m \{ y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \}$$

➔ Logistic regression cost function $J(\theta)$

$$J(\theta) = - \frac{1}{m} \sum_{i=1}^m \{ y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \}$$

➔ To fit parameters θ , find θ_{MLE} s.t. $\theta_{MLE} = \underset{\theta}{\min} J(\theta)$

Outline

- Classification
- Hypothesis representation
- Decision boundary
- Cost function
- Derivation of cost function via MLE
- Simplified cost function and gradient descent
- Advanced optimization
- Multi-class classification

Logistic Regression Cost Function

- Logistic regression cost function

$$J(\theta) = -\frac{1}{m} \left[\sum_{i=1}^m \{ y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \} \right]$$

- To fit parameters θ

- $\min_{\theta} J(\theta)$

- To make a prediction, given new x_{new} :

- $h_{\theta}(x_{new}) = \frac{1}{1 + \exp(-\theta^T x_{new})} = P(y = 1 | x_{new}; \theta)$

Gradient Descent

$$J(\theta) = - \frac{1}{m} \left[\sum_{i=1}^m \{ y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \} \right]$$

■ Want $\min_{\theta} J(\theta)$

■ Gradient descent

Repeat {

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

(simultaneously update for every θ_j)

■ $\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$

Gradient Descent

- $\frac{\partial}{\partial \theta_j} J(\theta) = \frac{\partial}{\partial \theta_j} \left(-\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right] \right)$
- $\frac{\partial}{\partial \theta_j} h_{\theta}(x^{(i)})$

Gradient Descent

- $\frac{\partial}{\partial \theta_j} J(\theta) = \frac{\partial}{\partial \theta_j} \left(-\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right] \right)$
- $\frac{\partial}{\partial \theta_j} \log h_{\theta}(x^{(i)})$

Gradient Descent

- $\frac{\partial}{\partial \theta_j} J(\theta) = \frac{\partial}{\partial \theta_j} \left(-\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right] \right)$
- $\frac{\partial}{\partial \theta_j} \log(1 - h_{\theta}(x^{(i)}))$

Gradient Descent

$$\blacksquare \frac{\partial}{\partial \theta_j} J(\theta) = \frac{\partial}{\partial \theta_j} \left(-\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right] \right)$$

$$\blacksquare \frac{\partial}{\partial \theta_j} J(\theta) =$$

Gradient Descent

- $\frac{\partial}{\partial \theta_j} J(\theta) = \frac{\partial}{\partial \theta_j} \left(-\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right] \right)$
- $\frac{\partial}{\partial \theta_j} h = \frac{1}{h} \left(x_j^{(i)} h(1 - h) \right)$
- $\frac{\partial}{\partial \theta_j} \log h_{\theta}(x^{(i)}) = \frac{1}{h} \frac{\partial}{\partial \theta_j} h = \frac{1}{h} \left(x_j^{(i)} h(1 - h) \right)$
- $\frac{\partial}{\partial \theta_j} \log(1 - h_{\theta}(x^{(i)})) = \frac{1}{1-h} \frac{\partial}{\partial \theta_j} (1 - h) = \frac{1}{1-h} \left(x_j^{(i)} h(1 - h) \right)$

$$\begin{aligned}
 \frac{\partial}{\partial \theta_j} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \left(y^{(i)} \frac{1}{h} x_j^{(i)} h(h - 1) + (1 - y^{(i)}) \frac{1}{1-h} \left(x_j^{(i)} h(h - 1) \right) \right) \\
 &= \frac{1}{m} \sum_{i=1}^m \left(y^{(i)} x_j^{(i)} (h - 1) - (1 - y^{(i)}) x_j^{(i)} h \right) \\
 &= \frac{1}{m} \sum_{i=1}^m \left(-y^{(i)} x_j^{(i)} + x_j^{(i)} h \right) = \frac{1}{m} \sum_{i=1}^m \left((h - y^{(i)}) x_j^{(i)} \right)
 \end{aligned}$$

Gradient Descent

$$J(\theta) = - \frac{1}{m} \left[\sum_{i=1}^m \{ y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \} \right]$$

■ Want $\min_{\theta} J(\theta)$

■ Gradient descent

Repeat {

$$\theta_j \leftarrow \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

} (simultaneously update for every θ_j)

■ Algorithm looks identical to linear regression!

■ Logistic regression: $h_{\theta}(z) = \frac{1}{1 + \exp(-\theta^T x)}$

■ Linear regression: $h_{\theta}(z) = \theta^T x$

Outline

- Classification
- Hypothesis representation
- Decision boundary
- Cost function
- Derivation of cost function via MLE
- Simplified cost function and gradient descent
- **Advanced optimization**
- Multi-class classification

Optimization Algorithm

- Cost function $J(\theta)$
- Want $\min_{\theta} J(\theta)$
- Given θ , we have code that can compute
 - $J(\theta)$
 - $\frac{\partial}{\partial \theta} J(\theta)$ (for $j = 0, 1, \dots, n$)
- Gradient descent
Repeat {

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}
(simultaneously update for every θ_j)

Optimization Algorithm

■ Given θ , we have code that can compute

■ $J(\theta)$

■ $\frac{\partial}{\partial \theta} J(\theta)$ (for $j = 0, 1, \dots, n$)

■ Optimization algorithms

■ Gradient descent

■ Conjugate gradient

■ BFGS

■ Broyden-Fletcher-Goldfarb-Shanno

■ L-BFGS

■ Limited memory - BFGS

■ Advantages

■ No need to manually pick α

■ Often faster than gradient descent

■ Disadvantages:

■ More complex

Outline

- Classification
- Hypothesis representation
- Decision boundary
- Cost function
- Derivation of cost function via MLE
- Simplified cost function and gradient descent
- Advanced optimization
- **Multi-class classification**

Multiclass Classification

■ Email foldering/tagging

■ Work($y = 1$), Friends ($y = 2$), Family ($y = 3$), Hobby ($y = 4$)



■ Medical diagrams

■ Not ill ($y = 1$), Cold ($y = 2$), Flu ($y = 3$),


■ Weather




■ Sunny ($y = 1$), Cloudy ($y = 2$), Rain ($y = 3$), Snow ($y = 4$),




Binary vs Multiclass Classification



One vs All (One vs Rest)



Class 1: 
 Class 2: 
 Class 3: 



$$h_{\theta_i}^{(i)}(x) = P(y = i | x; \theta_1, \theta_2, \theta_3) \quad (i = 1, 2, 3)$$

One vs All

- Train a logistic regression classifier $h_{\theta}^{(i)}(x)$ for each class i to predict the probability that $y = i$.
- On a new input x_{new} , to make a prediction, pick the class i that maximizes

$$\max_i h_{\theta}^{(i)}(x_{new})$$

Multi-Class Logistic Regression

■ $h_{\theta}(x)$

■ Estimated probability that $y = 1$ on input x

■ $p(y = 1|x; \theta)$

■ For binary classes

$$h_{\theta}(x) = \frac{1}{1 + \exp(-\theta^T x)} = \frac{\exp(\theta^T x)}{\boxed{1} + \boxed{\exp(\theta^T x)}}$$

Weight assigned to $y=0$

Weight assigned to $y=1$

Multi-Class Logistic Regression

- For K classes $\{1, 2, \dots, K\}$:

$$p(y = j | \mathbf{x}; \boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K) = \frac{\exp(\boldsymbol{\theta}_j^T \mathbf{x})}{\sum_{k=1}^K \exp(\boldsymbol{\theta}_k^T \mathbf{x})}$$

- Called the softmax function

- Given a sample vector \mathbf{x}

- Estimated probability for the j' th class

$$h_{\boldsymbol{\theta}}^{(j)}(\mathbf{x}) = p(y = j | \mathbf{x}; \boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K) = \frac{\exp(\boldsymbol{\theta}_j^T \mathbf{x})}{\sum_{k=1}^K \exp(\boldsymbol{\theta}_k^T \mathbf{x})} \quad (j = 1, \dots, K)$$


Softmax Function (Normalized Exponential Function)




$$\blacksquare p(y = j | \mathbf{x}; \boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K) = \frac{\exp(\boldsymbol{\theta}_j^T \mathbf{x})}{\sum_{k=1}^K \exp(\boldsymbol{\theta}_k^T \mathbf{x})}$$




- A generalization of the logistic regression that maps n -dim vector \mathbf{x} of real values to K -dim vector p of real values in the range $(0,1)$
- $\sum_{j=1}^K p(y = j | \mathbf{x}; \boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K) = 1$
- The output of softmax function can be used to
 - represent a categorical distribution
 - Probability distribution over K different possible outcomes

$$\blacksquare h_{\boldsymbol{\theta}}(\mathbf{x}) = \begin{bmatrix} h_{\boldsymbol{\theta}}^{(1)} \\ h_{\boldsymbol{\theta}}^{(2)} \\ \dots \\ h_{\boldsymbol{\theta}}^{(K)} \end{bmatrix} = \begin{bmatrix} p(y = 1 | \mathbf{x}; \boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K) \\ p(y = 2 | \mathbf{x}; \boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K) \\ \dots \\ p(y = K | \mathbf{x}; \boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K) \end{bmatrix}$$

Multi-Class Logistic Regression



Class 1: 
 Class 2: 
 Class 3: 



- Train a logistic regression classifier for each class j to predict the probability that $y = j$ with

$$h_{\theta}^{(j)}(\mathbf{x}) = p(y = j | \mathbf{x}; \theta_1, \dots, \theta_K) = \frac{\exp(\theta_j^T \mathbf{x})}{\sum_{k=1}^K \exp(\theta_k^T \mathbf{x})}$$

Implementing Multi-Class Logistic Regression

- As the model for class K,

- Use $h_{\theta}^{(j)}(\mathbf{x}) = p(y = j | \mathbf{x}; \theta_1, \dots, \theta_K) = \frac{\exp(\theta_j^T \mathbf{x})}{\sum_{k=1}^K \exp(\theta_k^T \mathbf{x})}$

- Gradient descent simultaneously updates all parameters for all models with the above $h_{\theta}^{(j)}(\mathbf{x})$

- Predict class label as the most probable label

$$\max_j h_{\theta}^{(j)}(\mathbf{x}) = \max_j p(y = j | \mathbf{x}; \theta_1, \dots, \theta_K)$$

References

- Andrew Ng, <https://www.coursera.org/learn/machine-learning>
- http://www.holehouse.org/mlclass/06_Logistic_Regression.html
- Eric Eaton, <https://www.seas.upenn.edu/~cis519>