

Machine Learning System Design

전 재 욱

**Embedded System 연구실
성균관대학교**

Outline

- Error Analysis
- Error metrics for skewed classes
- Data for machine learning

Outline

- Error Analysis
- Error metrics for skewed classes
- Data for machine learning

Recommended Approach

- Start with a simple algorithm that we can imp quickly.
 - Implement it and test it on our cross-validation data.

- Plot learning curves
to decide if more data, more features, etc. are likely to help.

- Error analysis
 - Manually examine the examples (in cross validation set)
that our algorithm made errors on.
 - See if we spot any systematic trend in what type of examples it is
making errors on.

Outline

- Error Analysis
- Error metrics for skewed classes
- Data for machine learning

Machine learning system design

■ Error metrics for skewed classes

Error Metrics for Skewed Analysis

- Example of skewed classes
 - (Number of examples in one class)
 << (Number of examples in the other)
 ➔ So standard error metrics are not so good

Cancer Classification Example

- Train logistic regression model $h_{\theta}(x)$

- ($y = 1$ if cancer, $y = 0$ otherwise)

- Find that you got 1% error on test set.

- (99% correct diagnoses)

- This looks pretty good.

- Only 0.5% of patients have cancer

```
function y = predictCancer(x)
    y = 0; %ignore x!
    return
```

→ 0.5% error

- Now, 1% error looks very bad.

- When the number of examples in one class is very small

- this is an example of skewed classes

Another Example

- Algorithm has 99.2% accuracy
 - Make a change, now get 99.5% accuracy
 - Does this really represent an improvement to the algorithm?

- Did we do something useful,
or did we just create something which predicts $y = 0$
more often
 - Get very low error, but classifier is still not great

Precision and Recall

- $y = 1$ in presence of rare class that we want to detect

		Actual Class	
		1	0
Predicted Class	1	True Positive	False Positive
	0	False Negative	True Negative

- Precision

- Of all patients where we predicted $y = 1$, what fraction actually has cancer?

- $$\frac{\text{True positives}}{\# \text{ of predicted positive}} = \frac{\text{True positive}}{\text{True pos} + \text{False Pos}}$$

- Recall

- Of all patients that actually have cancer, what fraction did we correctly detect as having cancer?

- $$\frac{\text{True positives}}{\# \text{ of actual positive}} = \frac{\text{True positive}}{\text{True pos} + \text{False neg}}$$

Precision and Recall

■ Precision

■ *How often does our algorithm cause a false alarm?*

■ Of all patients we predicted have cancer,
what fraction of them *actually* have cancer?

$$\text{■ } \frac{\text{True positives}}{\text{\# of predicted positive}} = \frac{\text{True positive}}{\text{True pos} + \text{False } \textcolor{red}{Pos}}$$

■ High precision is good (i.e. closer to 1)

■ We want a big number

➤ because we want *False* ***Positive*** to be as close to 0 as possible

Precision and Recall

Recall

- *How sensitive is our algorithm?*
- Of all patients in set that actually have cancer, what fraction did we correctly detect

$$\frac{\text{True positives}}{\# \text{ of actual positive}} = \frac{\text{True positive}}{\text{True pos} + \text{False neg}}$$

- High recall is good (i.e. closer to 1)
 - We want a big number
 - because we want *False negative* to be as close to 0 as possible

Precision and Recall

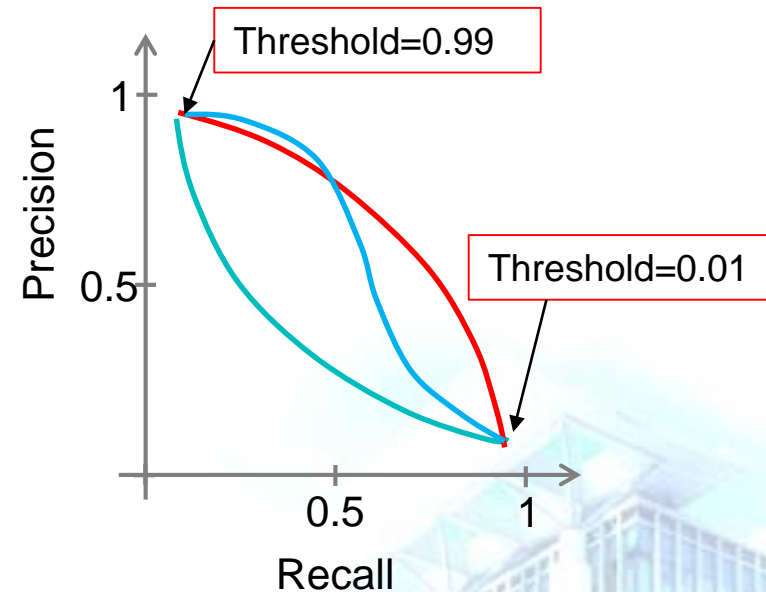
- By computing precision and recall
 - get a better sense of how an algorithm is doing
 - Means we are much more sure that an algorithm is good
- Typically, the presence of a rare class
 - is what we are trying to determine
 - e.g. positive (1) is the existence of the rare thing

Trading off Precision and Recall

- Logistic regression: $0 \leq h_{\theta}(x) \leq 1$
 - Predict 1 if $h_{\theta}(x) \geq 0.5$
 - Predict 0 if $h_{\theta}(x) < 0.5$
- Suppose we want to predict $y = 1$ (cancer) only if very confident.
 - Predict 1 if $h_{\theta}(x) \geq 0.8$
 - Higher precision, lower recall
 - Risk of false negatives
- Suppose we want to avoid missing too many cases of cancer (avoid false negatives).
 - Predict 1 if $h_{\theta}(x) \geq 0.3$
 - Higher recall, lower precision
 - Risk of false positives
- More generally:
 - Predict 1 if $h_{\theta}(x) \geq \text{thresholded}$

■ Precision = $\frac{\text{True positives}}{\# \text{ of predicted positive}}$

■ Recall = $\frac{\text{True positives}}{\# \text{ of actual positive}}$



Curve shape can be changed depending on classifier details

Trading off Precision and Recall

- How to compare precision/recall numbers?
 - Which algorithm is the best among the following three?

	Precision(P)	Recall(R)
Algo 1	0.5	0.4
Algo 2	0.7	0.1
Algo 3	0.02	1.0

Trading off Precision and Recall

■ How to compare precision/recall numbers?

■ Average: $\frac{P+R}{2}$

■ 0.45, 0.4, 0.51

➤ 0.51 is the best despite having a recall of 1 - i.e. predict $y=1$ for everything

■ NOT good

	Precision(P)	Recall(R)	Average
Algo 1	0.5	0.4	0.45
Algo 2	0.7	0.1	0.4
Algo 3	0.02	1.0	0.51

Trading off Precision and Recall

■ How to compare precision/recall numbers?

■ F_1 Score (F score): $2 \frac{PR}{P+R}$ ($\leftarrow \frac{2}{\frac{1}{p} + \frac{1}{R}}$)

■ F_1 score is like taking the average of precision and recall giving a higher weight to the lower value

■ $P=0$ or $R=0$

→ F_1 score = 0

■ $P=1$ **and** $R=1$

→ F_1 score = 1

	Precision(P)	Recall(R)	F Score
Algo 1	0.5	0.4	0.2222
Algo 2	0.7	0.1	0.0875
Algo 3	0.02	1.0	0.0196

Outline

- Error Analysis
- Error metrics for skewed classes
- Data for machine learning

Designing A High Accuracy Learning System

■ Classify btw confusable words

■ {to, two, too}, {then, than}

■ For breakfast I ate _____ eggs.

■ Algorithms

■ Perceptron (logistic regression)

■ Winnow

■ Like logistic regression

■ Used less now

■ Memory based

■ Used less now

■ Naive Bayes

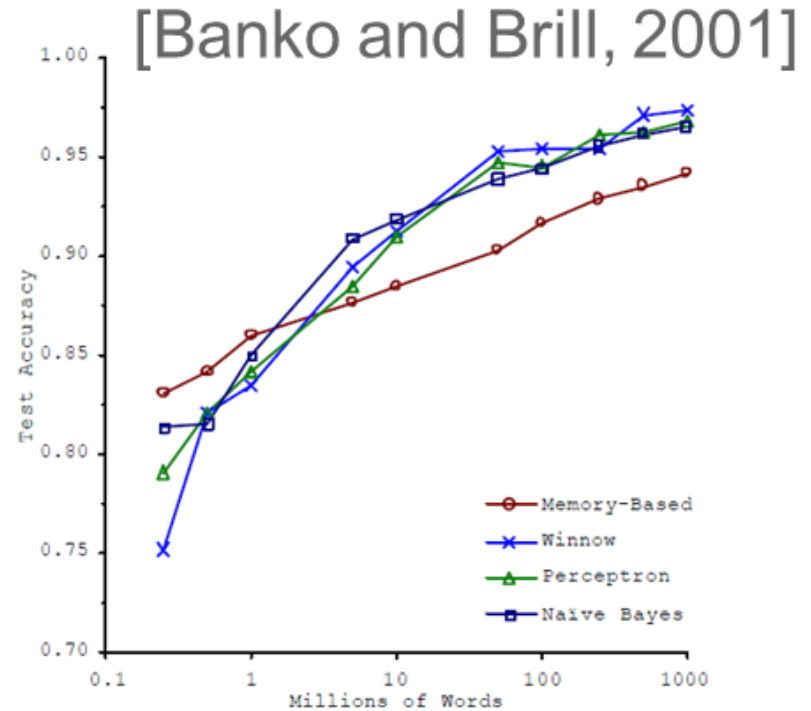


Figure 1. Learning Curves for Confusion Set Disambiguation

Designing A High Accuracy Learning System

What can we conclude

- Algorithms give remarkably similar performance
- As training set sizes increases, accuracy increases
- Take an algorithm, give it more data, should beat a "better" one with less data
- Shows that
 - Algorithm choice is pretty similar
 - More data helps




Figure 1. Learning Curves for Confusion Set Disambiguation

➔ It's not who has the best algorithm that wins.
It's who has the most data.

Large Data Rationale

- Assume feature $x \in R^{n+1}$ has sufficient information to predict y accurately.
 - Then more data may help
 - Example:
 - For breakfast I ate _____ eggs.
 - Counterexample:
 - Predict housing price from only size (feet²) and no other features.
- Useful test:
 - Given the input x , can a human expert confidently predict y ?

Large Data Rationale

- Use a learning algorithm with many parameters
 - e.g. logistic regression/linear regression with many features
 - neural network with many hidden units
- A powerful learning algorithm with many parameters which can fit complex functions
 - Low bias algorithm
 - Little systemic bias in their description – flexible
 - $J_{train}(\theta)$ will be small
- Use a very large training set
 - Unlikely to overfit
 - $J_{train}(\theta) \approx J_{test}(\theta)$

$J_{test}(\theta)$ will be small

Large Data Rationale

- An algorithm with having both low bias and low variance
 - Use complex algorithm
 - For low bias
 - Use large training set
 - For low variance

References

- Andrew Ng, <https://www.coursera.org/learn/machine-learning>
- http://www.holehouse.org/mlclass/11_Machine_Learning_System_Design.html
- M. Banko and E. Brill, "Scaling to Very Very Large Corpora for Natural Language Disambiguation," 2001.